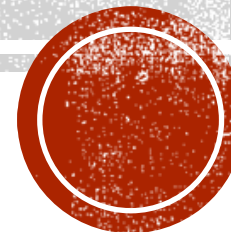



Local Search



LOCAL SEARCH AND OPTIMIZATION

- Previous methods: path to goal is a solution to a problem
 - systematic exploration of search space.
- Today approaches: a state is a solution to problem
 - for some problems path is irrelevant.
 - E.g., 8-queens
- Different algorithms can be used
 - Depth First Branch and Bound
 - Local search





LOCAL SEARCH AND OPTIMIZATION

Local search

- Keep track of single current state
- Move only to neighboring states
- Ignore paths

Advantages:

- Use very little memory
- Can often find reasonable solutions in large or infinite (continuous) state spaces.

“Pure optimization” problems

- All states have an objective function
- Goal is to find state with max (or min) objective value
- Does not quite fit into path-cost/goal-state formulation
- Local search can do quite well on these problems.



HILL CLIMBING

“a loop that continuously moves towards increasing value”

- terminates when a peak is reached
- Aka greedy local search

Value can be either

- Objective function value
- Heuristic function value (minimized)


Hill climbing does not look ahead of the immediate neighbors

Can randomly choose among the set of best successors

- if multiple have the best value

“climbing Mount Everest in a thick fog with amnesia”





HILL CLIMBING

It starts from an initial solution and evolves that single solution into a mostly better and better solution.

Example n -queens (here $n=4$)

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Successor function

move a single queen to another square in the same column.

h = number of pairs of queens that are attacking each other

- $h = 6$ for initial state




HILL CLIMBING

It starts from an initial solution and evolves that single solution into a mostly better and better solution.

Example n -queens (here $n=4$)

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Successor function

move a single queen to another square in the same column.

h = number of pairs of queens that are attacking each other

- $h = 6$ for initial state



HILL CLIMBING

It starts from an initial solution and evolves that single solution into a mostly better and better solution.

Example n -queens (here $n=4$)


Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Successor function

move a single queen to another square in the same column.

h = number of pairs of queens that are attacking each other

- $h = 6$ for initial state



HILL CLIMBING

It starts from an initial solution and evolves that single solution into a mostly better and better solution.

Example n -queens (here $n=4$)

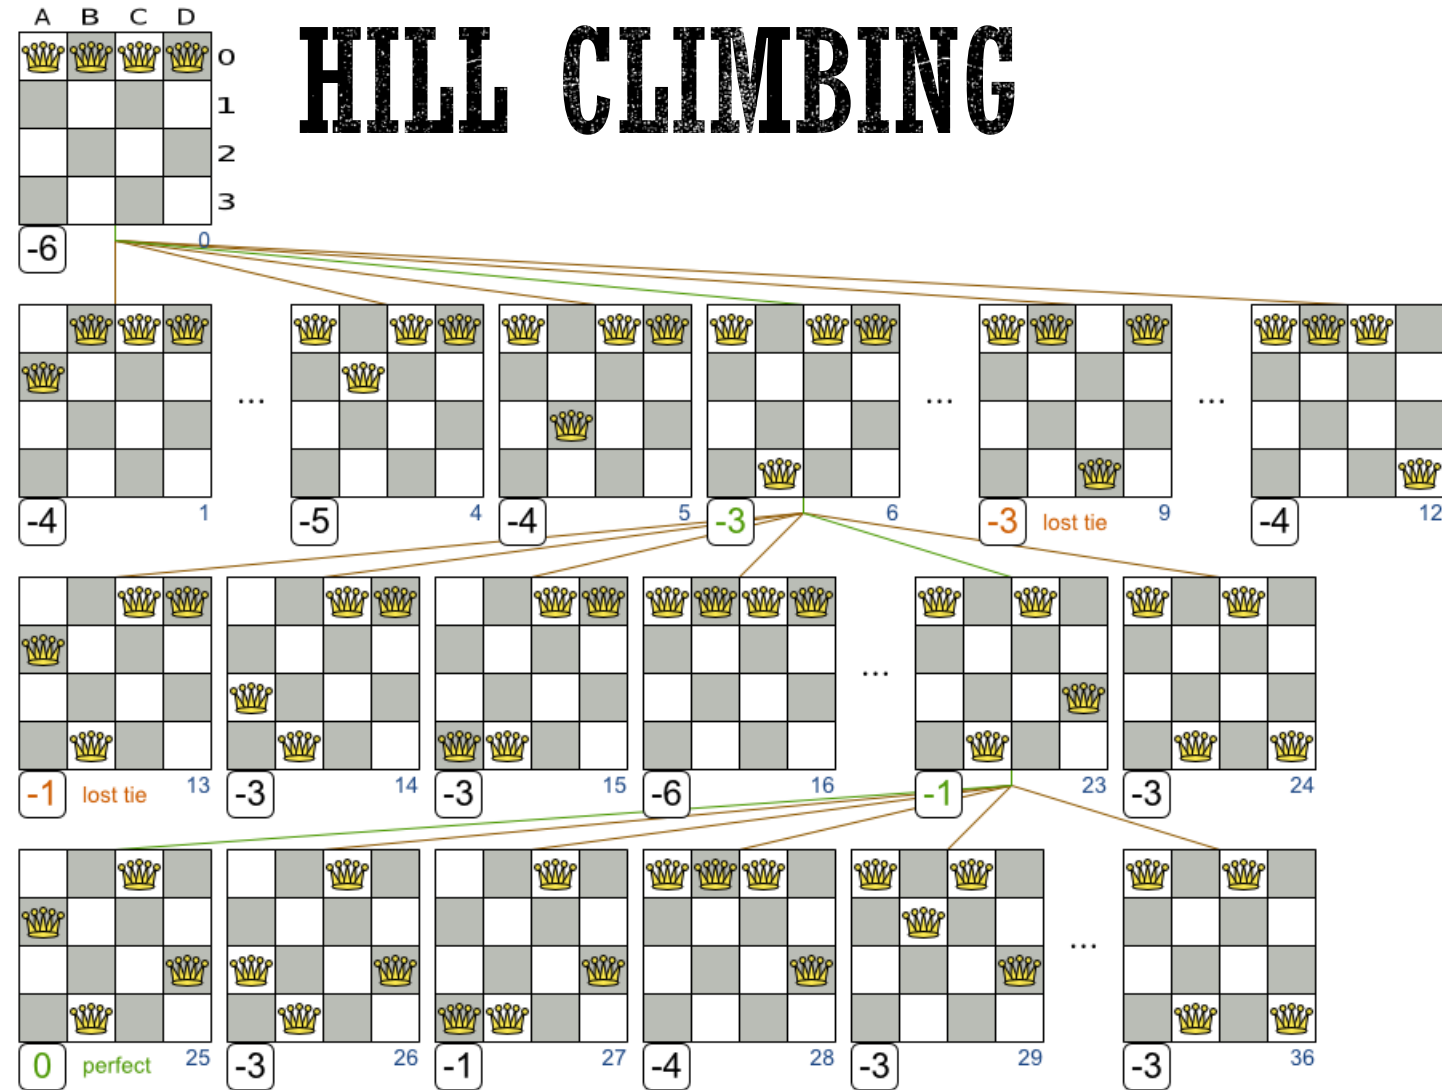
Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Successor function

move a single queen to another square in the same column.

h = number of pairs of queens that are attacking each other

- $h = 6$ for initial state



HILL CLIMBING

function HILL-CLIMBING(*problem*) **return** a state that is a local maximum

input: *problem*, a problem

local variables: *current*, a node.

neighbor, a node.

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

loop do

neighbor \leftarrow a highest valued successor of *current*

if VALUE [*neighbor*] \leq VALUE[*current*] **then return** STATE[*current*]

current \leftarrow *neighbor*

min version will reverse inequalities and look for lowest valued successor




HILL CLIMBING

Example of 8 queens Problem

- Randomly generated 8-queens starting states...
- 14% the time it solves the problem
- 86% of the time it get stuck at a local minimum

It takes only 4 steps on average when it succeeds


- And 3 on average when it gets stuck
- (for a state space with $8^8 \approx 17$ million states)



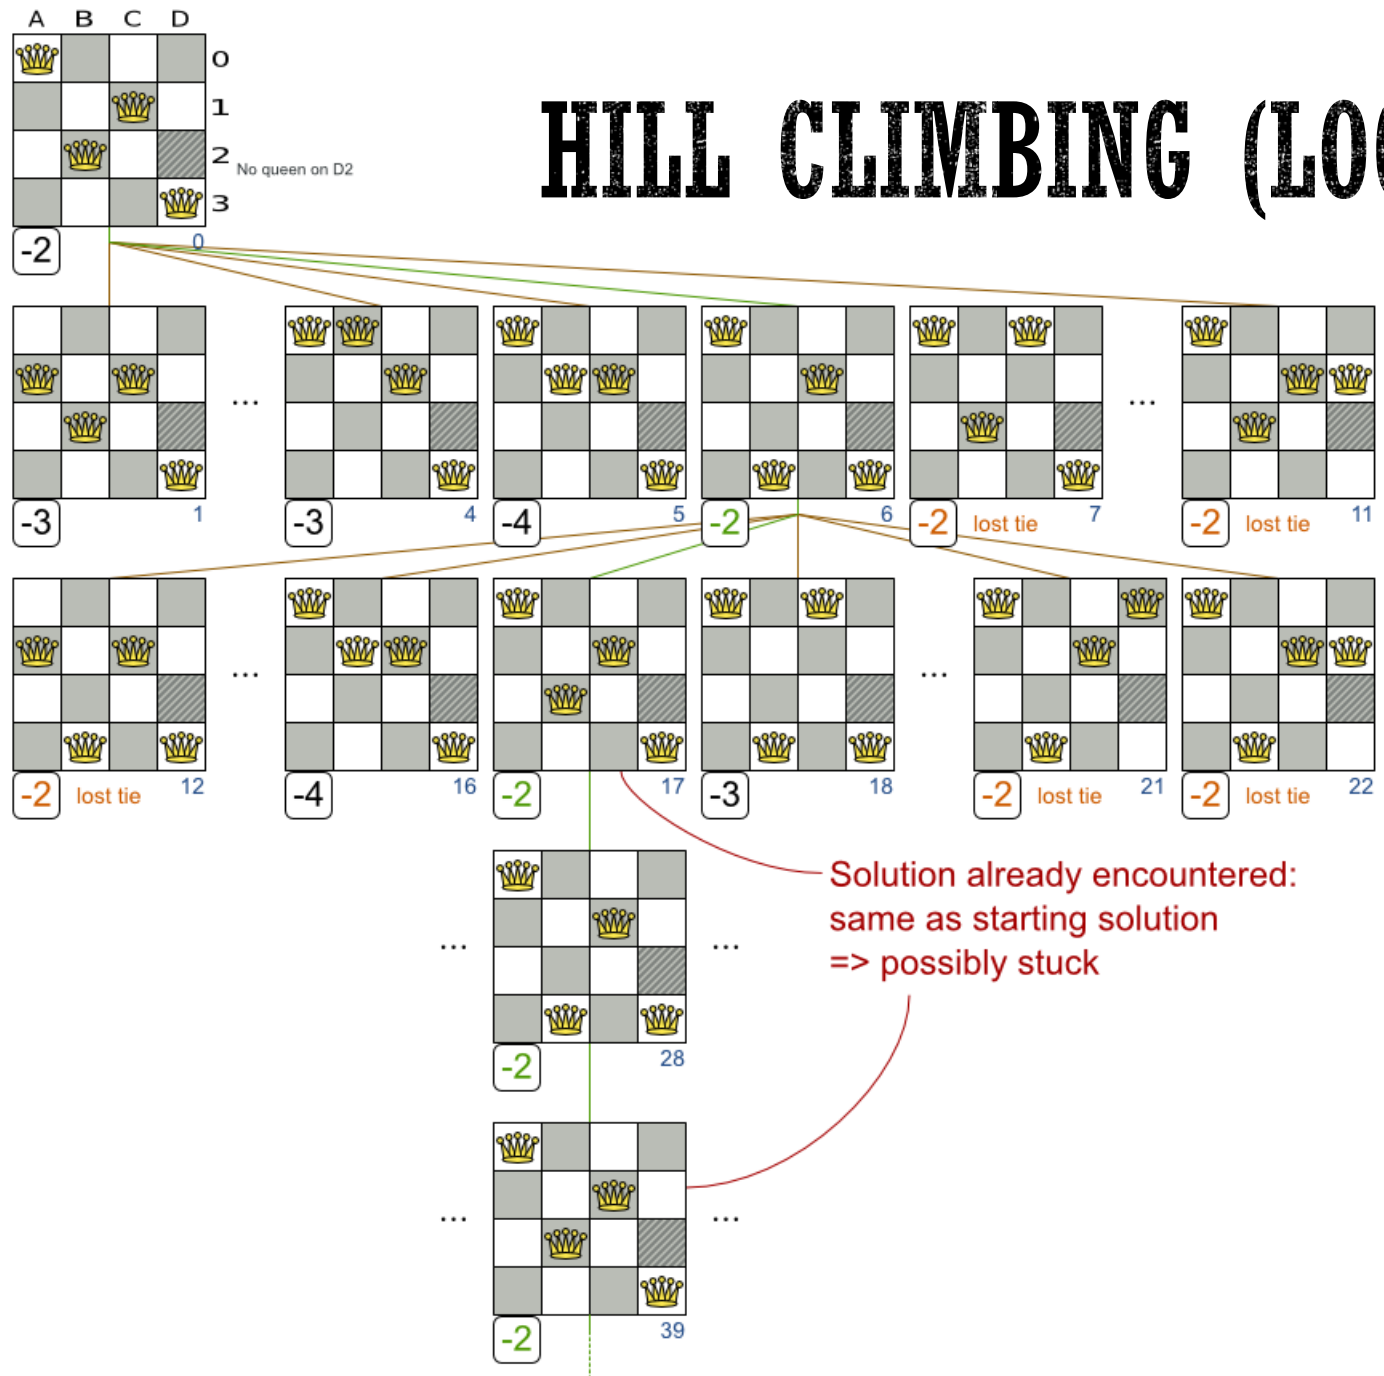
Is this a solution?



HILL CLIMBING



HILL CLIMBING (LOCAL OPTIMA PROBLEM)



HILL-CLIMBING: STOCHASTIC VARIATIONS

Stochastic hill-climbing

- Random selection among the uphill moves.
- The selection probability can vary with the steepness of the uphill move.

To avoid getting stuck in local minima

- Random-walk hill-climbing
- Random-restart hill-climbing
- Hill-climbing with both




HILL-CLIMBING WITH RANDOM RESTARTS

If at first you don't succeed, try, try again!

- Different variations
 - For each restart: run until termination vs. run for a fixed time
 - Run a fixed number of restarts or run indefinitely

***** an Interesting Approach *****



SIMULATED ANNEALING

Simulated Annealing = physics inspired twist on random walk

- Basic ideas:
 - like hill-climbing identify the quality of the local improvements
 - instead of picking the best move, pick one randomly
 - say the change in objective function is δ
 - if δ is positive, then move to that state
 - otherwise:
 - move to this state with probability proportional to δ
 - thus: worse moves (very large negative δ) are executed less often
 - however, there is always a chance of escaping from local maxima
 - over time, make it less likely to accept locally bad moves
 - (Can also make the size of the move random as well, i.e., allow “large” steps in state space)






PHYSICAL INTERPRETATION OF SIMULATED ANNEALING

- **A Physical Analogy:**
 - imagine letting a ball roll downhill on the function surface
 - this is like hill-climbing (for minimization)
 - now imagine shaking the surface, while the ball rolls, gradually reducing the amount of shaking
 - this is like simulated annealing
- Annealing = physical process of cooling a liquid or metal until particles achieve a certain frozen crystal state
 - simulated annealing:
 - free variables are like particles
 - seek “low energy” (high quality) configuration
 - slowly reducing temp. T with particles moving around randomly



Temperature
decreases
for each step

Step 0		
t	Δ	max 
2.0	≥ 0	any
	-1	0.61
	-2	0.37
	-3	0.22
	-4	0.14



Simulated Annealing (Time Gradient aware)

N queens ($n = 4$, starting Temperature = 2)






Temperature
decreases
for each step

Step 0

t	Δ	max
2.0	≥ 0	any
	-1	0.61
	-2	0.37
	-3	0.22
	-4	0.14

Step 1

t	Δ	max
1.6	≥ 0	any
	-1	0.54
	-2	0.29
	-3	0.15
	-4	0.08



Simulated Annealing (Time Gradient aware)

N queens (n = 4, starting Temperature = 2)




Temperature
decreases
for each step

Simulated Annealing (Time Gradient aware)

N queens (n = 4, starting Temperature = 2)


Step 0

t	Δ	max
2.0	≥ 0	any
	-1	0.61
	-2	0.37
	-3	0.22
	-4	0.14




Step 1

t	Δ	max
1.6	≥ 0	any
	-1	0.54
	-2	0.29
	-3	0.15
	-4	0.08



Step 2

t	Δ	max
1.2	≥ 0	any
	-1	0.43
	-2	0.19
	-3	0.08
	-4	0.04




Step 3

t	Δ	max
0.8	≥ 0	any
	-1	0.29
	-2	0.08
	-3	0.02
	-4	0.01



Step 4

t	Δ	max
0.4	≥ 0	any
	-1	0.08
	-2	0.01
	-3	0.00
	-4	0.00



SIMULATED ANNEALING

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a solution state

input: *problem*, a problem

schedule, a mapping from time to temperature

local variables: *current*, a node.

next, a node.

T, a “temperature” controlling the prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E / T}$

high T: probability of “locally bad” move is higher

low T: probability of “locally bad” move is lower

typically, T is decreased as the algorithm runs longer

i.e., there is a “temperature schedule”



GENETIC ALGORITHMS

Twist on Local Search: successor is generated by combining two parent states

A state is represented as a string over a finite alphabet (e.g. binary)

- 8-queens
 - State = position of 8 queens each in a column

Start with k randomly generated states (**population**)

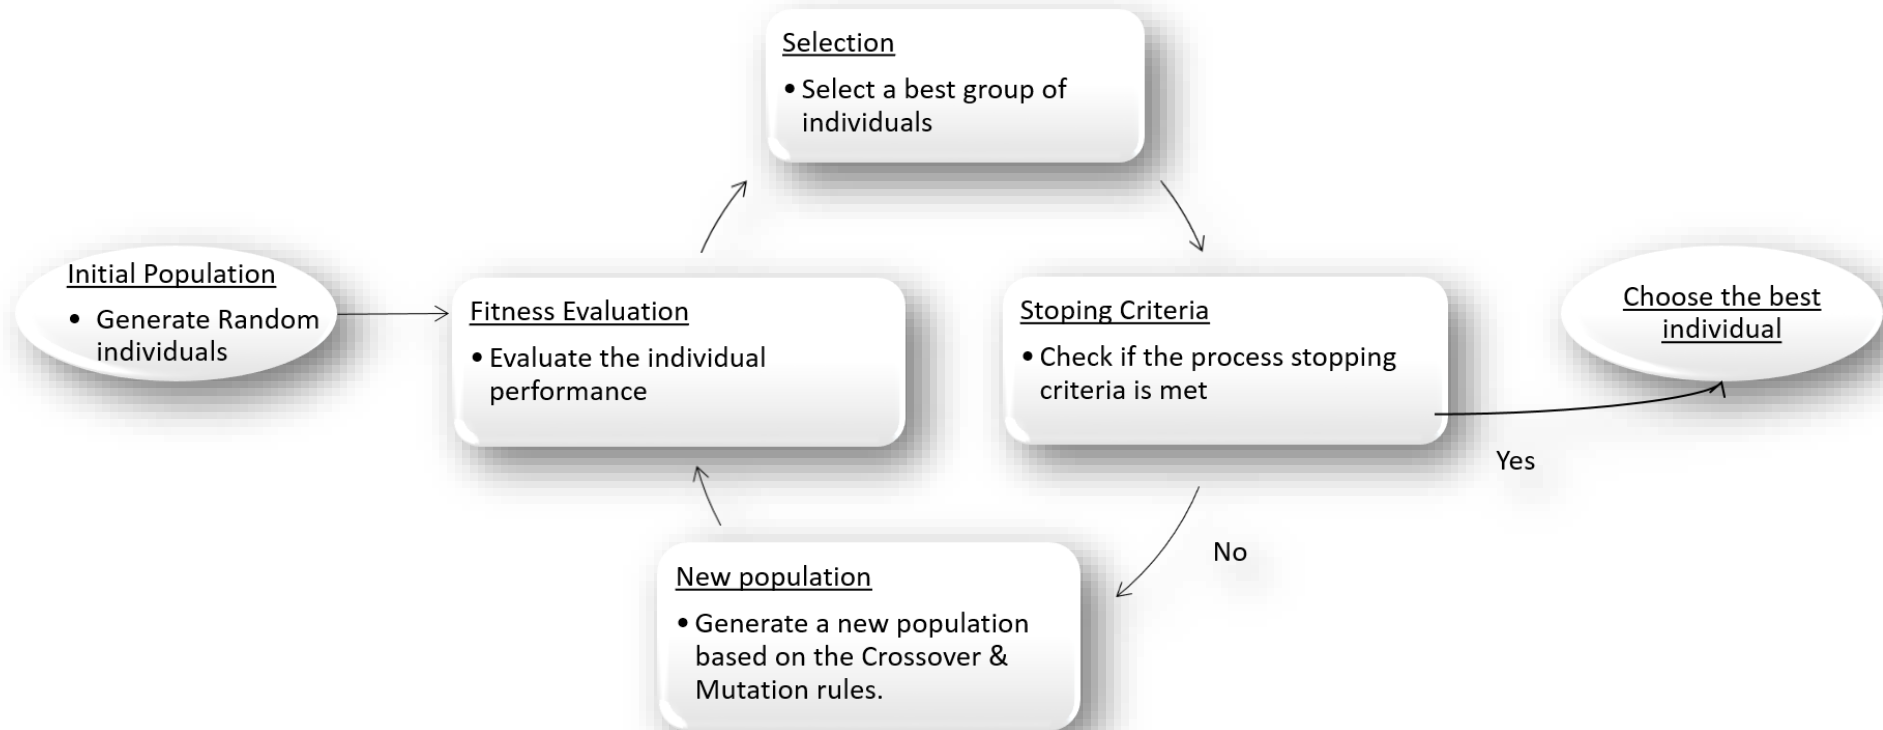
Evaluation function (**fitness function**):

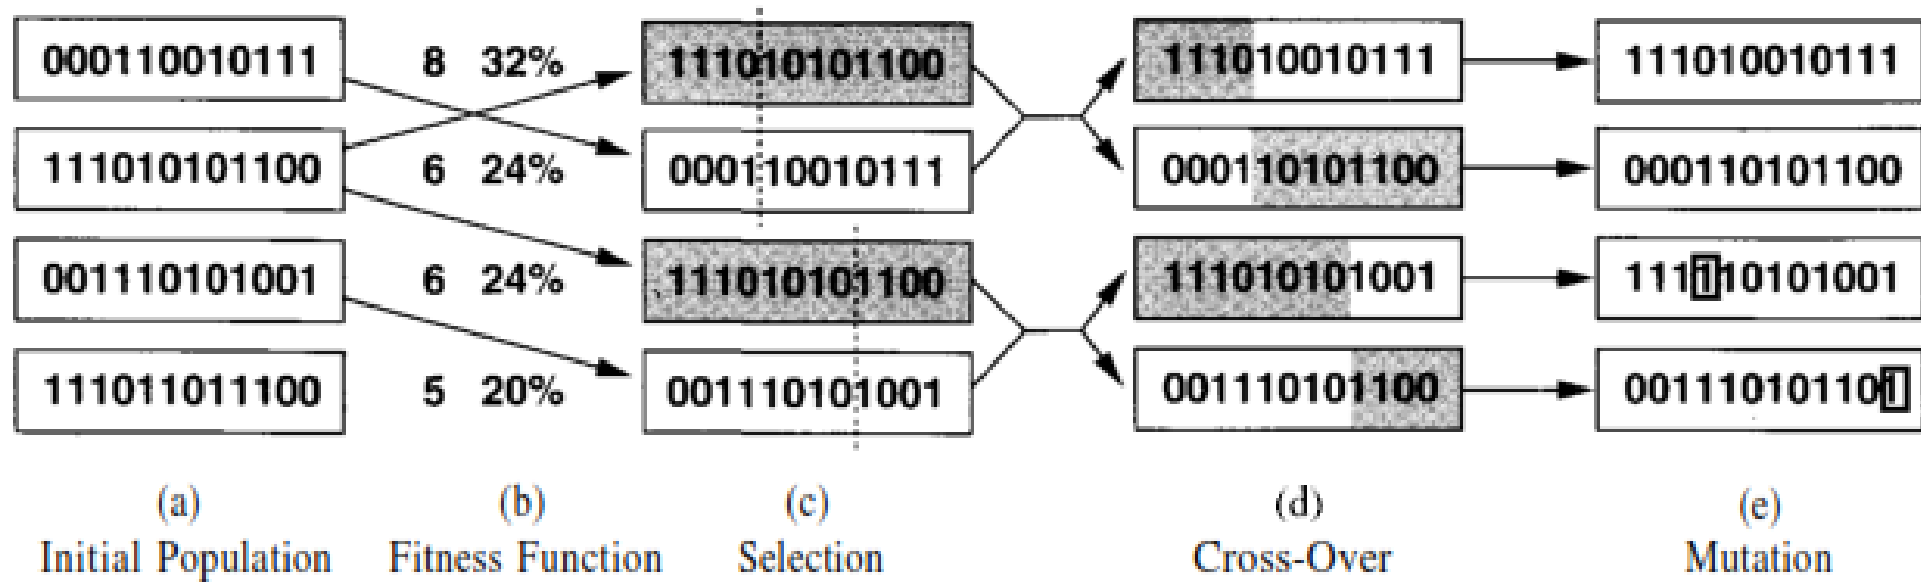
- Higher values for better states.
- Opposite to heuristic function, e.g., # non-attacking pairs in 8-queens

Produce the next generation of states by “simulated evolution”

- Random selection
- Crossover
- Random mutation












Source : Peter norvig genetic algorithm illustration





Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)

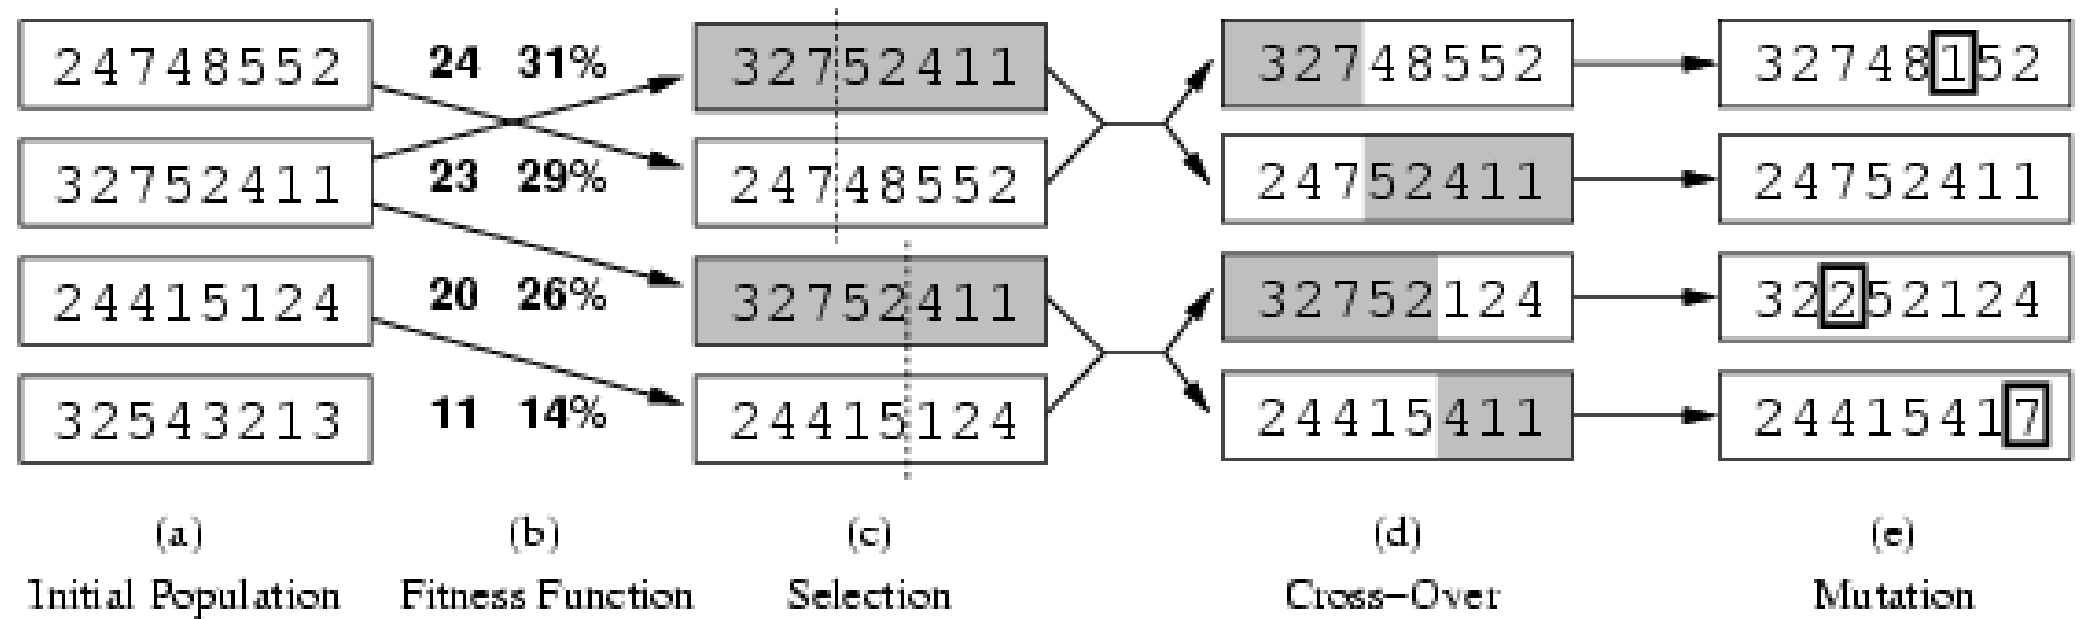
[7,5,1,4,6,3,2,0]

[0,5,1,4,6,3,7,2]

- [1,1,1,2,3,3,4,5]
- [2,4,7,4,8,5,5,2]
- [3,2,7,5,2,4,1,1]

- [2,4,4,1,5,1,2,4]
- [3,2,5,4,3,2,1,3]
- [0,1,2,6,6,7,3,1]





COMMENTS ON GENETIC ALGORITHMS

- Positive points
 - Random exploration can find solutions that local search can't
 - (via crossover primarily)
 - Appealing connection to human evolution
 - “neural” networks, and “genetic” algorithms are **metaphors!**
- Negative points
 - Large number of “tunable” parameters
 - Difficult to replicate performance from one problem to another

